

## CS 444/544 Spring 2020 Lab 2

Lab 2 is due by 11:59pm on Tuesday, March 31st. Don't plan on any extensions, Lab 3 will also need a substantial amount of time.

Submit your lab in Learn. It will appear as an “exam” with two fields. In one you’ll copy and paste the plaintext phrase you decrypted. It will be in English and will be unmistakable if you crack the code. You can escape tab or newline characters, but there shouldn’t be any. In the other you’ll upload a single gzipped tar ball containing all your source code that is specific to your attack. Your source code need not run to be graded, but I reserve the right to call you into my office to demo it for me and it will be required to run on your machine for that.

Lab 2 is worth 100 points, all or nothing. If you fail to include your source code as a tar ball, or your plaintext does not match the plaintext that was encrypted in the PCAP specifically assigned to you as a student, then you will get 0 points. I won't grade your source code for style, *etc.*, and I won't attempt to run it. The source code is a fallback if I suspect any students of cheating.

You are expected to do your own work. From writing the source code to carrying out the attack, it should be an individual effort carried out only by you. Any instance of not doing your own work will be considered cheating. You are encouraged to discuss the assignment with your classmates at a high level. Exchanging publicly available source code that existed before the assignment was assigned, and thoughts about approaches to specific problems is okay. Accessing or seeing in any way source code written for the related assignment written by a Fall 2018 student of this class is forbidden. It is also okay to go through attack examples together at a high level (*e.g.*, pencil and paper), but do not share any concrete info (such as a network capture or source code) about any specific attack. As a reminder of the course policy, if you cheat on any assignment in this class including this assignment (cheating includes, but is not limited to, representing somebody else's work as your own or fabricating files or text to make it look like you completed the assignment) you will receive an F in the class. Every student's session key and plaintext/ciphertext are unique, and you should not extract the ciphertext or attempt to recover the plaintext of any other student. You may use the Python code (written by Meisam Navaki and Nick Aase) provided as part of the tar balls for the assignment however you like without restriction (*e.g.*, you can run your own server locally for testing, you can write your own script that calls `rsa-aes-client.py`, you can borrow from the provided code or turn it into a library, use it as a starting point for your own code, or whatever).

Your mission, should you choose to accept it (April 17<sup>th</sup> is the last day to drop without the Dean’s permission) is to extract the ciphertext corresponding to the port that will be assigned to you, and then use a padding oracle attack to recover the plaintext that corresponds to your ciphertext. The server is listening on ports 10000-10100 on a server at IP address 64.106.39.33, but you **must only use the port that is assigned to you**. You are not limited in the number of queries you can submit to the server, but please be respectful of the resources. I recommend making no more than 1 query every 5 seconds to the server on your port. If you need more than 512 connections for any attempt, then you're doing it wrong. Any kind of denial-of-service attack or other attack on the server itself is NOT authorized. If you'd like to try an alternative method to complete the assignment, ask me first and I might give you authorization to try to break into the server (but only if you show me something specific that’s worth trying).

A unique student number, that corresponds to your port, will be assigned to you *via* Learn. Do not extract ciphertexts that weren't assigned to you or connect to ports that weren't assigned to you. You

should always connect to port 10000 plus your student number. For example, if you are student number 55, you will always only connect to port 10055 on yuba. Connecting to ports not assigned to you will be considered cheating (because you're interfering with other students' ability to complete the assignment and/or trying to complete their assignment for them, probably). Extracting ciphertexts, or asking any oracle for information about the corresponding plaintext, for any ciphertext not assigned to you will also be considered cheating (because there is no reason to do this unless you are doing someone else's assignment). Honest mistakes will be forgiven, but still be very careful to only work with your own ciphertext and your own port. Do not share your assigned ciphertext or plaintext with anyone else, or tell them your port number. I reserve the right to record all network traffic to and from yuba to catch and build a case against any cheaters.

A tarball will be provided, and you probably already found it if you're reading this. It will help you to set up your own server for local testing, and includes the public key that corresponds to the private key the server uses. There will also be a packet capture, only one TCP session of which you should extract (based on your student number). We will discuss how to carry out the attack in class, and I'll send out a sample chapter and/or a paper about the attack.

Basically, every PCAP is a connection the TA made to the server on a given port. Each connection used a randomly generated 256-bit AES key to encrypt a message. That AES key was encrypted using the server's 2048-bit RSA private key, and the message and encrypted key were sent to the server. The server is a capitalization service, and you have the source code for the client and server. You also have the RSA public key used to encrypt the AES session key. The only thing you lack is the RSA private key, which only exists on the server. Your mission is to recover the message. You'll recover the AES session key in the process of doing so, but you will probably never know the RSA private key. In fact, you don't need it to recover the plaintext. You'll do a padding oracle attack to recover the RSA plaintext as it is leaked one bit at a time by the server, and the RSA plaintext is the AES session key.

To find out your student number, check Learn. There will be an assignment (which won't actually count towards your grade), and your grade between 0 and 100 on that assignment will tell you your student number.

We tried to filter out inappropriate plaintexts, but may have missed one. Please know that plaintexts were randomly assigned to students.

Notes:

- The ciphertext assigned to you is the one that you'll extract from the PCAP file for your port number. So if your student number is 55, your port is 10055.
- There are easy ways to extract the ciphertext with Wireshark. You can do C Arrays and just copy the client part, or do hex and use vim to fix it up.
- I **strongly** recommend that you set up your own server on the loopback interface to develop your attack source code, and then once you've tested it against your own encrypted messages *then and only then* start making queries against the server (on the port assigned to you).
- Feel free to email me with questions, but I also intent to set up a forum where you're more likely to get timely help because I'll check it regularly (more so than my email) and the TA and other students will also be able to help you there.
- Feel free to use ports 10080 through 10099 if you need to (*e.g.*, if your port freezes up and doesn't come back), but use them at your own risk since other students may be using them. Do not use port 10000, that port is reserved for the TA.
- I recommend using exception handling to keep trying to test bits, and test them again if there is

some kind of server error. The attack can *in theory* be done with 256 queries (or maybe even less), but it may be best to make sure a bit is correct before trying to move on, and sometimes you'll attempt to connect to the server and get a disconnection error, so you'll end up making more than 256 queries and making more attempts than you have successes at querying (probably).

- **Start early.** The attack seems simple and coding up the logic for it isn't too bad, but getting the socket stuff right and debugging your code when there's encryption involved can lead to some pretty hairy bugs. If you get stuck, ask for help. But be sure to **start early** and get past those bugs sooner rather than later.